

# Artificial Intelligence



## Constraint Satisfaction Problems (Part 2)

CS 444 – Spring 2021

Dr. Kevin Molloy

Department of Computer Science

James Madison University

Much of this lecture is taken from  
Dan Klein and Pieter Abbeel AI class at UC Berkeley

# Today

---

- Review of A\* Heuristics for PA 1
- Continue with Constraint Satisfaction problems
- Arc consistency AC-3 examples
- Problem Structure
- Min conflicts

# Learning Objectives

---

- Apply the AC-3 to maintain arc consistency (MAC)
- Investigate the problem structure of CSPs for identify more efficient solutions using cutset conditioning and tree decomposition
- Apply min-conflicts algorithm and by able to code it to solve CSPs. Characterize the min-conflicts algorithm (runtime, completeness, etc).

# Student Heuristic Presentation

---

- Alex Marasco – Finding/visit all the corners heuristics
- Garrett Christian -- Eat all the dots heuristic

# CSP problems

- CSPs:
  - Variables
  - Domains
  - Constraints
    - Implicit (provide code to compute)
    - Explicit (provide a list of the legal tuples)
  - Types:
    - Unary (one variable)
    - Binary (two variables)
    - N-ary (n variables)
- Goals:
  - **In this class:** find any solution
  - Also: find all, find best, etc.






Figure from Berkley AI

# Backtracking Search

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
    return RECURSIVE-BACKTRACKING({ }, csp)
function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
    if assignment is complete then return assignment
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
        if value is consistent with assignment given CONSTRAINTS[csp] then
            add {var = value} to assignment
            result  $\leftarrow$  RECURSIVE-BACKTRACKING(assignment, csp)
            if result  $\neq$  failure then return result
            remove {var = value} from assignment
    return failure
```

# Backtracking Example






# Improving Backtracking

- General-purpose ideas give huge gains in speed
  - ... but it's all still **NP-hard**
- **Filtering:** Can we detect inevitable failure early?
- **Ordering:**
  - Which **variable** should be assigned next? (MRV)
  - In what order should its **values** be tried? (LCV)
- **Structure:** Can we exploit the problem structure?



# Consistency of a Single Arc

- An arc  $X \rightarrow Y$  is **consistent** iff for every  $x$  in the tail there is *some*  $y$  in the head which could be assigned without violating a constraint




*Delete from the tail!*

- Forward checking: Enforcing consistency of arcs pointing to each new assignment

# Arc Consistency of an Entire CSP

- A simple form of propagation makes sure **all** arcs are consistent:



- Important: If X loses a value, neighbors of X need to be rechecked!
- Arc consistency detects failure **earlier** than forward checking
- Can be run as a preprocessor or after each assignment
- What's the downside of enforcing arc consistency?

*Remember: Delete  
from the tail!*

# Enforcing Arc Consistency in a CSP

```
function AC-3( csp ) returns the CSP, possibly with reduced domains
    inputs: csp, a binary CSP with variables  $\{X_1, X_2, \dots, X_n\}$ 
    local variables: queue, a queue of arcs, initially all the arcs in csp

    while queue is not empty do
         $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\textit{queue})$ 
        if REMOVE-INCONSISTENT-VALUES( $X_i, X_j$ ) then
            for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
                add  $(X_k, X_i)$  to queue




---


function REMOVE-INCONSISTENT-VALUES(  $X_i, X_j$  ) returns true iff succeeds
    removed  $\leftarrow \text{false}$ 
    for each  $x$  in DOMAIN[ $X_i$ ] do
        if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x, y)$  to satisfy the constraint  $X_i \leftrightarrow X_j$ 
            then delete  $x$  from DOMAIN[ $X_i$ ]; removed  $\leftarrow \text{true}$ 
    return removed
```

- Runtime:  $O(n^2d^3)$ , can be reduced to  $O(n^2d^2)$
- ... but detecting all possible future problems is NP-hard – why?

# Limitations of Arc Consistency

- After enforcing arc consistency:
  - Can have one solution left
  - Can have multiple solutions left
  - Can have no solutions left (and not know it)
- Arc consistency still runs inside a backtracking search!



*What went  
wrong here?*

# K-Consistency

- Increasing degrees of consistency
  - 1-Consistency (Node Consistency): Each single node's domain has a value which meets that node's unary constraints
  - 2-Consistency (Arc Consistency): For each pair of nodes, any consistent assignment to one can be extended to the other
  - K-Consistency: For each k nodes, any consistent assignment to k-1 can be extended to the k<sup>th</sup> node.
- Higher k more expensive to compute
- (You need to know the k=2 case: arc consistency)





Figure from Berkley AI

# Strong K-Consistency

---

- Strong k-consistency: also k-1, k-2, ... 1 consistent
- Claim: strong n-consistency means we can solve without backtracking!
- Why?
  - Choose any assignment to any variable
  - Choose a new variable
  - By 2-consistency, there is a choice consistent with the first
  - Choose a new variable
  - By 3-consistency, there is a choice consistent with the first 2
  - ...
- Lots of middle ground between arc consistency and n-consistency! (e.g. k=3, called path consistency)

# Intelligent Backtracking




Variable assignment order: {Q, NSW, V, T, SA, WA, NT}

Partial Assignment: {Q = red, NSW=green, V=blue, T=red}

What does normal backtracking do when it tries to assignment **SA** a color?

# Intelligent Backtracking



Variable assignment order: {Q, NSW, V, T, SA, WA, NT}


Partial Assignment: {Q = red, NSW=green, V=blue, T=red}

What does normal backtracking do when it tries to assignment **SA** a color?

- It tries all 3 colors. None of these work. So backtrack.
- Change the color of T and try SA again
- Still no assignment works for SA. So backtrack.
- Etc.

How can we make this better?

# Intelligent Backtracking – Back jumping



Variable assignment order: {Q, NSW, V, T, SA, WA, NT}

Partial Assignment: {Q = red, NSW=green, V=blue, T=red}


Idea: Jump to a variable that is causing a problem.

Define a **conflict set**, which is built as we evaluate a variable. So, for SA, we check:

- Can't use red, Q is added to the conflict set for SA.
- Can't use green, NSW is added to the conflict set for SA.
- Can't use blue, V is added to the conflict set for SA.

Backtrack to at least one of these variables so we have a chance of correcting the issue.


# Tree-Structured CSPs




- Theorem: if the constraint graph has no loops, the CSP can be solved in  $O(n d^2)$  time
  - Compare to general CSPs, where worst-case time is  $O(d^n)$
- This property also applies to probabilistic reasoning (later): an example of the relation between syntactic restrictions and the complexity of reasoning

# Tree-Structured CSPs

- Algorithm for tree-structured CSPs:
  - Order: Choose a root variable, order variables so that parents precede children




- Remove backward: For  $i = n : 2$ , apply RemoveInconsistent( $\text{Parent}(X_i), X_i$ )
- Assign forward: For  $i = 1 : n$ , assign  $X_i$  consistently with  $\text{Parent}(X_i)$
- Runtime:  $O(n d^2)$  (why?)




# Tree-Structured CSPs

- Claim 1: After backward pass, all root-to-leaf arcs are consistent
- Proof: Each  $X \rightarrow Y$  was made consistent at one point and  $Y$ 's domain could not have been reduced thereafter (because  $Y$ 's children were processed before  $Y$ )



- Claim 2: If root-to-leaf arcs are consistent, forward assignment will not backtrack
- Proof: Induction on position
- Why doesn't this algorithm work with cycles in the constraint graph?
- Note: we'll see this basic idea again with Bayes' nets

# Nearly Tree-Structured CSPs



- **Conditioning:** instantiate a variable, prune its neighbors' domains
- **Cutset conditioning:** instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree
- Cutset size  $c$  gives runtime  $O( (d^c) (n-c) d^2 )$ , very fast for small  $c$


# Cutset Conditioning

Choose a cutset

Instantiate the cutset  
(all possible ways)


Compute residual CSP  
for each assignment

Solve the residual CSPs  
(tree structured)





# Cutset Quiz

- Find the smallest cutset for the graph below.




# Tree Decomposition

- Idea: create a tree-structured graph of mega-variables
- Each mega-variable encodes part of the original CSP
- Subproblems overlap to ensure consistent solutions



# Iterative Algorithms for CSPs

- Local search methods typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
  - Take an assignment with unsatisfied constraints
  - Operators *reassign* variable values
  - No fringe! Live on the edge.
- Algorithm:



While not solved:

Variable selection: randomly select any conflicted variable

Value selection: min-conflicts heuristic:

Choose a value that violates the fewest constraints

I.e., hill climb with  $h(n) = \text{total number of violated constraints}$

# Performance of Different CSP Algorithms

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV	Min-Conflicts
USA (4 color)	(> 1,000,000)				
$n$ -Queens	(> 40,000,000)				

# Performance of Different CSP Algorithms

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV	Min-Conflicts
USA (4 color)	(> 1,000,000)	(> 1,000,000)			
$n$ -Queens	(> 40,000,000)	13,500,000			

# Performance of Different CSP Algorithms

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV	Min-Conflicts
USA (4 color)	(> 1,000,000)	(> 1,000,000)	2,000		
$n$ -Queens	(> 40,000,000)	13,500,000	(> 40,000,000)		

# Performance of Different CSP Algorithms

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV	Min-Conflicts
USA (4 color)	(> 1,000,000)	(> 1,000,000)	2,000	60	
$n$ -Queens	(> 40,000,000)	13,500,000	(> 40,000,000)	817,000	

# Performance of Different CSP Algorithms

Problem	Backtracking	BT+MRV	Forward Checking	FC+MRV	Min-Conflicts
USA (4 color)	(> 1,000,000)	(> 1,000,000)	2,000	60	64
$n$ -Queens	(> 40,000,000)	13,500,000	(> 40,000,000)	817,000	4,000

# Don't Make Things too Complicated

- Given random initial state, can solve n-queens in almost constant time for arbitrary n with high probability (e.g., n = 10,000,000)!
- The same appears to be true for any randomly-generated CSP *except* in a narrow range of the ratio

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$





Figure from Berkley AI

# CSP Summary

- CSPs are a special kind of search problem:
  - States are partial assignments
  - Goal test defined by constraints
- Basic solution: backtracking search
- Speed-ups:
  - Ordering
  - Filtering
  - Structure
- Iterative min-conflicts is often effective in practice

